

Operating Systems – Exercise 1

System Calls, Basic I/O

Submission & General Guidelines

- Submission deadline is **11/4/2018, 23:55** Moodle server time
- This exercise must run properly on the provided virtual machine
- Submit your answers in the course website only as single ex1-YOUR_ID.zip (e.g. ex1-012345678.zip), containing only:
 - ex1.pdf
 - ex1.c
- Place your name and ID at the top of every source file, as well as in the PDF with the answers.
- No late submission will be accepted.
- Do not submit handwritten work.
- Please provide concise answers, but make sure to explain them.
- Write clean code(readable, documented, consistent, etc...)

Part 1 (36 points)

In this question we will implement a C application that copies a file. The application should receive the source file path, target file path and buffer size (in bytes) as command line arguments. By default, the application does not overwrite an existing file, unless the -f option was specified.

1. Read the manual page for the following System Calls i.e. execute: `man 2 read`
 - a. `READ(2)`
 - b. `WRITE(2)`
 - c. `OPEN(2)`

Carefully read about the following option flags:
`O_RDONLY`, `O_WRONLY`, `O_CREAT`, `O_TRUNC`, `O_EXCL`
 - d. `CLOSE(2)`
2. Read the manual page for the following C Library Calls
 - a. `PERROR(3)`
 - b. `PRINTF(3)`
 - c. `EXIT(3)`
3. Complete the application ex1.c (missing parts are marked with `//TODO`)

Execution output examples:

```
$ ./ex1
Invalid number of arguments
Usage:
    ex1 [-f] SOURCE DEST BUFFER_SIZE

$ ./ex1 /etc/passwd
/tmp/passwd Invalid number of
arguments Usage:
    ex1 [-f] SOURCE DEST BUFFER_SIZE

$ ./ex1 /etc/passwd /tmp/passwd 4096
File /etc/passwd was copied to /tmp/passwd

$ ./ex1 /etc/passwd /tmp/passwd 4096
Unable to open destination file for writing: File exists

$ ./ex1 -f /etc/passwd /tmp/passwd 4096
File /etc/passwd was copied to /tmp/passwd

$ ./ex1 -f /tmp/passwd /etc/passwd 4096
Unable to open destination file for writing: Permission denied

$ ./ex1 /etc/password /tmp/passwd 4096
Unable to open source file for reading: No such file or directory
```

Guidelines

- Use the manuals. Chapters 2 & 3 are your friends
- Make sure to always close the files you are using
- Always check system calls return value for errors. ALWAYS.
- You are not allowed to use any external / C-Library code that copies files, if you are not sure if you are allowed to use something, ask in the forum.
- Your program must finish executing with `EXIT_SUCCESS` only when there were no errors (otherwise it must finish executing with `EXIT_FAILURE` after printing a helpful message).
- **Hint:** the 4 system calls and 3 libc calls mentioned above are all **you** need to implement this part of the exercise successfully.

Part 2 (24 points)

We will now examine the performance of our program from part 1.

1. Create a 5MB file using the following command:

```
dd if=/dev/zero of=/tmp/test.5mb bs=1M count=5
```

2. Run your program on this file, preceded with the **time(1)** command, using the following buffer sizes: 128, 256, 512, 1024, 2048, 4096, 8192, 16384, 32768, 65536

```
$ time ./ex1 -f /tmp/test.5mb /tmp/test.5mb_dest 65536
File /tmp/test.5mb was copied to /tmp/test.5mb_dest

real    0m0.004s
user    0m0.000s
sys     0m0.000s
```

3. Draw a graph (using Google Sheets or Microsoft Excel) with a single series. The graph should show the **real** time each run takes (in milliseconds) [Y axis], as a function of the buffer size [X axis]
4. Explain the graph. Why isn't it a straight line, parallel to the X axis?
5. Hypothetically, if we change the **copy_file** in part 1 to print out to the screen a message each time we read **buffer_size** bytes, will the running time change significantly?
 - **IMPORTANT:** Do not make this change. Submission with this change will result in 0 points for this question!

Part 3 (40 points)

For each of the statements below, state whether **it is true or false and explain concisely** (two lines at most):

1. In Virtualization, the hypervisor emulates the hardware for the guest OS and therefore it does not need to use system calls (This question assumes the existence of a host OS).
2. In Virtualization, guest operating systems can access hardware components at the same speed as the host operating system does.
3. Usually, if a program uses more system calls it will run slower.
4. It is up to the developer of a program to decide whether his program will run in User mode or Kernel mode.

5. The hypervisor, being a software that manages other VMs, is always running on Kernel mode (This question assumes the existence of a host OS).
6. Applications such as a web browser (e.g. Chrome/Explorer) are running in user mode and therefore are not allowed to invoke system calls.
7. A simple program which prints "Welcome to the Operating Systems course!" to the monitor must run entirely in Kernel mode in order to access the monitor.
8. Pressing the left button of the mouse will cause an interrupt while pressing the right button will not.
9. External hardware devices may access the operating system by using System Calls.
10. A developer can choose to limit the number of interrupts that occur during the execution of his program.